

Data Analysis With Pandas I & II

October 7, 2024

1 Data Analysis with Pandas

Michael Fudge 2024 Edition

1.0.1 Table of Contents:

0. What is Pandas? What is Data Analysis?
1. **Part One:**
 1. Importing Pandas and numpy
 2. Basic Elements of Pandas: Series
 3. Basic Elements of Pandas: DataFrame
 4. Accessing elements with loc and iloc
 5. Basic DataFrame operations: info(), count(), describe(), head(), tail(), sample()
 6. Selecting Rows and Columns: selectors and boolean indexes
2. **Part Two:**
 1. Iterating over a DataFrame or series
 2. Loading JSON with Pandas, json_normalize()
 3. Simple Web scraping with Pandas
 4. Advanced Dataframe Operations: concat(), merge()
 5. Generating Columns, Lambdas and the apply() function
 6. DataFrames as output

1.1 What is Pandas?

Pandas is a Python library for data analysis. The homepage of Pandas is <http://pandas.pydata.org/>.

Pandas takes the pain and suffering out of data analysis by doing a lot of the work for you. Pandas is so elegant at times, people feel it's "indistinguishable from magic." In this lesson we will try to our best to explain what Pandas is doing "behind the curtain" and expose the magic behind Pandas.

The Pandas library gets its name not from the furry Asian animal, but from the words "Panel Data".

1.2 What is data analysis?

Data Analysis is the process of systematically applying statistical and/or logical techniques to describe and illustrate, condense and recap, and evaluate data https://ori.hhs.gov/education/products/n_illinois_u/datamanagement/datopic.html.

The goals of data analysis are to:

1. discover useful information,
2. provide insights,
3. suggest conclusions, and
4. support decision-making.

1.2.1 Some examples

Most data analysis problems start with a fundamental question (but they don't always have to). Examples:

- Do students who study in groups perform better on examinations than those who study alone?
- What role (if any at all) does weather play in consumer shopping habits?
- What types of passengers were most likely to survive the Titanic?
- Among American Football teams who "go for it" on 4th downs, what is their win percentage?
- Are we stocking products in warehouses closest to the customers who purchase them?

1.2.2 Data analysis is "information storytelling"

Don't think of data analysis as crunching numbers and churning out graphs.

I like to think of data analysis as "information storytelling." Unlike slapping a chart on a powerpoint slide, which might be the *result* of a data analysis, it's a full disclosure of the process:

1. helping the reader understand your methodology (how you acquired and prepared the data)
2. sharing your complete analysis (including things that didn't work)
3. providing a narrative as to what the results mean, and most importantly
4. providing an honest and accurate analysis.

1.2.3 Why perform a data analysis with Code?

You **could** do this in Excel, but Code (Python) and Jupyter Notebook offer several advantages:

- It is easier to automate and update the process later on, since its code.
- You can intermix code with a narrative / explanations in the Jupyter notebook
- It can integrate with a variety of services and systems (because its code!).
- What you create is reproduceable. Someone else can run your code and see your results. This provides stronger evidence to support your analysis as opposed to just a chart on a powerpoint slide!

Data analysis is a skill every information professional should have, and this is a primary reason we teach programming in the iSchool. When you can code it makes performing a data analysis that much better!

1.3 Part One

1.3.1 1.A Import pandas

Before you can use **pandas** you must import it into Python. When we import Pandas it is customary to use the alias **pd** so that we don't have to type **pandas** in our code. For example instead of typing `pandas.DataFrame()` we can save ourselves a few keystrokes with `pd.DataFrame()`

np.nan In most cases we will need **np.nan** from the **numpy** module, too. This is a special value which indicates the absence of a value. (nan means “Not a Number”) Since these values are absent they are not included in numerical calculations. This is important for getting the “math” right, as we will see later on.

```
[1]: import pandas as pd
import numpy as np

print(np.nan)
```

nan

1.3.2 1.B The Basic Elements of Pandas: Series

Series The **series** is the most basic structure in Pandas. It is simply a named-list of values. Every series has a data type, expressed as a numpy data type. Common data types you will see in this course are **int** or **float64** for numbers, and **Object** for string types.

The following example creates a series called **grades** we use Python’s named arguments feature to set the data, name and data type (**dtype**) arguments. Some of the grades are missing so we use **np.nan** in those cases.

```
[2]: grades = pd.Series(data = [100,80,100,90,np.nan,70], name = "Grade")
grades
```

```
[2]: 0    100.0
1     80.0
2    100.0
3     90.0
4      NaN
5     70.0
Name: Grade, dtype: float64
```

Index What’s with those numbers to the left of the grades? That’s the **index**, and every series has one. An index is an ordered list of values. It’s how Pandas accesses specific values in a Series (or as we will see in a bit... a DataFrame)

The index of the series works a lot like the index in a list or a string. For example, this code prints the first and last grades in the series.

```
[3]: print("first grade:", grades[0])
print("last grade:", grades[5])
```

```
first grade: 100.0
last grade: 70.0
```

Series Aggregate Functions The beauty of Pandas is that for most operations, you won’t even need a loop! For example, we derive the **min()**, **max()**, **mean()**, **sum()** and **count()** of non **np.nan** values in a series without a for loop!!!

We do this by calling those method functions on the series itself, for example:

```
[4]: print("Highest grade:", grades.max())
      print("Average grade:", grades.mean())
      print("lowest grade:", grades.min())
      print("Sum of grades:", grades.sum())
      print("Count of grades", grades.count())
```

```
Highest grade: 100.0
Average grade: 88.0
lowest grade: 70.0
Sum of grades: 440.0
Count of grades 5
```

Other Series Functions You can also get value counts for a series and also deduplicate values in a series. For example consider the following series:

```
[5]: votes = pd.Series(data=[ 'y', 'y', 'y', 'n', 'y', np.nan, 'n', 'n', 'y'], name="Vote")
      votes
```

```
[5]: 0      y
      1      y
      2      y
      3      n
      4      y
      5    NaN
      6      n
      7      n
      8      y
      Name: Vote, dtype: object
```

We use the `unique()` method function to return only the non-duplicate values from the series.

The `value_counts()` method function adds up values, creating a new series where the index is the value and the value is the count.

```
[6]: print("deduplicate the votes:", votes.unique())
      print("counts by value:", votes.value_counts())
```

```
deduplicate the votes: ['y' 'n' nan]
counts by value: Vote
y      5
n      3
Name: count, dtype: int64
```

Where did all these functions come from? They're methods (object functions) associated with the Series object. If you `dir(grades)` you can see them!

If you want to learn more, here's the official documentation: <https://pandas.pydata.org/pandas-docs/stable/reference/series.html>

1.3.3 1.C Basic Elements of Pandas: DataFrame

The pandas **DataFrame** is a table representation of data. It is the primary use case for pandas itself. A dataframe is simply a collection of **Series** that share a common **Index**. I like to think of the DataFrame as a programmable spreadsheet. It has rows and columns which can be accessed and manipulated with Python.

DataFrame is the most common Pandas data structure. As you'll see its expressive and versatile, making it an essential tool in data analysis.

To make a **DataFrame** we must create a Python dict of Series, or list . We use the series name as the key and the series itself as the value.

This example creates a DataFrame from two series of Student names and Grade-Point-Averages (GPA's), using Series:

```
[7]: names = pd.Series( data = ['Allen', 'Bob', 'Chris', 'Dave', 'Ed', 'Frank', 'Gus'])
      gpas = pd.Series( data = [4.0, np.nan, 3.4, 2.8, 2.5, 3.8, 3.0])
      years = pd.Series( data = ['So', 'Fr', 'Fr', 'Jr', 'Sr', 'Sr', 'Fr'])
      series_dict = { 'Name': names, 'GPA': gpas, 'Year' : years } # dict of
      ↪Series, keys are the series names
      students = pd.DataFrame( series_dict )
      students
```

```
[7]:
```

	Name	GPA	Year
0	Allen	4.0	So
1	Bob	NaN	Fr
2	Chris	3.4	Fr
3	Dave	2.8	Jr
4	Ed	2.5	Sr
5	Frank	3.8	Sr
6	Gus	3.0	Fr

Here's the same code but we make the dataframe from simple Python lists. The end result is the same, as the DataFrame is always constructed from Series. In this case the `pd.DataFrame()` method created the three **Series** for us.

```
[8]: data_dict = {
      'Name': ['Allen', 'Bob', 'Chris', 'Dave', 'Ed', 'Frank', 'Gus'],
      'GPA': [4.0, np.nan, 3.4, 2.8, 2.5, 3.8, 3.0],
      'Year' : ['So', 'Fr', 'Fr', 'Jr', 'Sr', 'Sr', 'Fr'] }
      students = pd.DataFrame( data_dict )
      students
```

```
[8]:
```

	Name	GPA	Year
0	Allen	4.0	So
1	Bob	NaN	Fr
2	Chris	3.4	Fr
3	Dave	2.8	Jr
4	Ed	2.5	Sr

```
5 Frank 3.8 Sr
6 Gus 3.0 Fr
```

Getting a series from the dataframe You can access a series from the dataframe using the series name. For example this gets the lowest GPA by using the `min()` function on the `students['GPA']` series:

```
[9]: print("Lowest GPA:", students['GPA'].min())
```

```
Lowest GPA: 2.5
```

Getting the column names You can use the `columns` property to return an iterable of the columns in the dataframe.

```
[10]: print("Columns in the students data frame are: ", students.columns)
```

```
Columns in the students data frame are: Index(['Name', 'GPA', 'Year'],
dtype='object')
```

1.3.4 1.D Accessing elements with loc and iloc

The `loc` and `iloc` properties allow you to slice the dataframe. `loc` uses the index and column names, while `iloc` uses ordinal positions starting at zero.

Here are some examples:

```
[11]: # Examples using loc
print("loc: Get the name of the student at index 3: ", students.loc[3, 'Name'])
print("loc: Get the Year of the last student: ", students.loc[6, 'Year'])

# Same examples using iloc
print("iloc: Get the name of the student at index 3: ", students.iloc[3, 0])
print("iloc: Get the Year of the last student: ", students.iloc[-1, 2])
```

```
loc: Get the name of the student at index 3: Dave
loc: Get the Year of the last student: Fr
iloc: Get the name of the student at index 3: Dave
iloc: Get the Year of the last student: Fr
```

You can also use `loc` and `iloc` to retrieve a subset dataframe. For example, here are the names and GPA's of the first 4 students:

```
[12]: students.loc[0:3, 'Name':'GPA'] #uses the index and column names
```

```
[12]:      Name  GPA
0  Allen  4.0
1    Bob  NaN
2  Chris  3.4
3   Dave  2.8
```

```
[13]: students.iloc[0:4, 0:2] # uses relative places, zero-based.
```

```
[13]:      Name  GPA
0  Allen  4.0
1    Bob  NaN
2  Chris  3.4
3   Dave  2.8
```

```
[16]: students[ students.GPA.isna()]
```

```
[16]:      Name  GPA Year
1    Bob  NaN   Fr
```

1.3.5 1.E Basic DataFrame operations

In most cases, you will not construct a **DataFrame** as we have done in the examples so far. Instead you will load data from a file into a **DataFrame**. **pandas** supports a wide variety of file formats, which we will demonstrate in a later section, for now let's load a dataset of customers from the internet and use dataframe operations to explore the data.

This example read in a text file in **CSV** format into the **DataFrame**. There are 30 customers.

```
[14]: customers = pd.read_csv('https://raw.githubusercontent.com/mafudge/datasets/
↳master/customers/customers.csv')
len(customers)
```

```
[14]: 30
```

Getting information about your DataFrame First thing you are likely to do with a loaded dataframe it get some basic information about the data. These methods will help:

- **info()** provide names of columns, counts of non-null values in each columns, and data types.
- **describe()** for each numerical column provide some basic statistics (min, max, mean, and quartiles).

```
[15]: customers.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   First                 30 non-null    object
1   Last                  30 non-null    object
2   Email                 30 non-null    object
3   Gender                30 non-null    object
4   Last IP Address       30 non-null    object
5   City                  30 non-null    object
```

```

6   State          30 non-null   object
7   Total Orders   30 non-null   int64
8   Total Purchased 30 non-null   int64
9   Months Customer 30 non-null   int64
dtypes: int64(3), object(7)
memory usage: 2.5+ KB

```

```
[16]: customers.describe()
```

```

[16]:      Total Orders  Total Purchased  Months Customer
count      30.000000      30.000000      30.000000
mean        4.633333     1193.166667     15.133333
std         4.295012     1685.396507     15.734288
min         0.000000         0.000000         0.000000
25%         1.000000         68.750000         2.000000
50%         3.500000        500.000000         8.000000
75%         7.750000       1168.750000        27.750000
max        14.000000       6090.000000        42.000000

```

Viewing Data in your DataFrame Most likely the data in your DataFrame will be too large to view on the screen. There are methods to help with this:

- `head(n=5)` view the FIRST `n` rows in the dataframe (defaults to 5)
- `tail(n=5)` view the LAST `n` rows in the dataframe (defaults to 5)
- `sample(n=1)` view a random `n` rows from the dataframe (defaults to 1)

```
[17]: customers.sample(n=3)
```

```

[17]:      First  Last      Email Gender Last IP Address      City \
24   Sal    Ladd  sladd@superrito.com      M   23.112.202.16  Rochester
7    Candi  Cayne  ccayne@rhyta.com      F    24.39.14.15   Portland
18   Mike  Rofone  mrofone@dayrep.com      M    23.224.160.4   Cheyenne

      State  Total Orders  Total Purchased  Months Customer
24   NY          14          594          10
7    ME           1          620           2
18   WY           0           0           0

```

```
[18]: customers.head(n=3)
```

```

[18]:      First  Last      Email Gender Last IP Address      City State \
0     Al  Fresco  afresco@dayrep.com      M   74.111.18.161  Syracuse  NY
1   Abby    Kuss   akuss@rhyta.com      F   23.80.125.101   Phoenix  AZ
2  Arial  Photo  aphoto@dayrep.com      F    24.0.14.56    Newark  NJ

      Total Orders  Total Purchased  Months Customer
0              1              45              1
1              1              25              2

```


2 1 680 1

```
[19]: customers.tail(n=3)
```

```
[19]:
```

	First	Last	Email	Gender	Last IP Address	City \
27	Tally	Itupp	titupp@superrito.com	F	24.38.114.105	Sea Cliff
28	Tim	Pani	tpani@superrito.com	M	23.84.132.226	Buffalo
29	Victor	Rhee	vrhee@einrot.com	M	23.112.232.160	Green Bay

	State	Total Orders	Total Purchased	Months Customer
27	NY	11	380	42
28	NY	0	0	1
29	WI	0	0	2

1.3.6 1.F Selecting Rows and Columns: selectors and boolean indexes

In this final section of part one, we learn how to select rows and columns from the DataFrame using the selector operator []

```
[20]: students
```

```
[20]:
```

	Name	GPA	Year
0	Allen	4.0	So
1	Bob	NaN	Fr
2	Chris	3.4	Fr
3	Dave	2.8	Jr
4	Ed	2.5	Sr
5	Frank	3.8	Sr
6	Gus	3.0	Fr

Selecting Columns To select columns from the dataframe, include a list of columns you would like to include inside the selector.

For example this includes the **Year** and **Name** columns from the **students** DataFrame.

```
[21]: students[['Year', 'Name']]
```

```
[21]:
```

	Year	Name
0	So	Allen
1	Fr	Bob
2	Fr	Chris
3	Jr	Dave
4	Sr	Ed
5	Sr	Frank
6	Fr	Gus

Here another example where we include the name, email city and state from the **customers** dataframe. Since there are a lot of customers we use **head()** to show the first 5.

```
[22]: customers.columns
```

```
[22]: Index(['First', 'Last', 'Email', 'Gender', 'Last IP Address', 'City', 'State',  
        'Total Orders', 'Total Purchased', 'Months Customer'],  
        dtype='object')
```

```
[23]: customers[['First', 'Last', 'Email', 'City', 'State']].head(n=5)
```

```
[23]:
```

	First	Last	Email	City	State
0	Al	Fresco	afresco@dayrep.com	Syracuse	NY
1	Abby	Kuss	akuss@rhyta.com	Phoenix	AZ
2	Arial	Photo	aphoto@dayrep.com	Newark	NJ
3	Bette	Alott	balott@rhyta.com	Raleigh	NC
4	Barb	Barion	bbarion@superrito.com	Dallas	TX

Selecting A single column: DataFrame or Series? When selecting a single column, you have a decision to make: do you want a **Series** or a **DataFrame**?

- For a **Series** just include the column name in the selector: `'Email'`
- For a **DataFrame** include the column name as a list in the selector: `['Email']`

Again, in this example, `sample()` will be used to keep the size down!

```
[24]: # Series  
customers.sample(5)['Email']
```

```
[24]: 6      bmelator@einrot.com  
14     jpoole@dayrep.com  
9      crha@einrot.com  
18     mrofone@dayrep.com  
15     lhmeehom@einrot.com  
Name: Email, dtype: object
```

```
[25]: # DataFrame  
customers.sample(5)[['Email']]
```

```
[25]:
```

	Email
26	tanott@rhyta.com
10	ddelyons@dayrep.com
17	mmelator@rhyta.com
4	bbarion@superrito.com
0	afresco@dayrep.com

Selecting Rows: Boolean Index In pandas, a **boolean index** is a **Series** of type `bool` based on the result of some boolean expression.

We can then use the pandas selector `[]` to apply the boolean index to the **DataFrame** returning only rows where the boolean index is **True**

```
[26]: students
```

```
[26]:
```

	Name	GPA	Year
0	Allen	4.0	So
1	Bob	NaN	Fr
2	Chris	3.4	Fr
3	Dave	2.8	Jr
4	Ed	2.5	Sr
5	Frank	3.8	Sr
6	Gus	3.0	Fr

For example this boolean index evaluates to `True` when the value in the `'Year'` series is equal to `'Fr'`.

```
[27]: students['Year'] == 'Fr'
```

```
[27]:
```

0	False
1	True
2	True
3	False
4	False
5	False
6	True

Name: Year, dtype: bool

To filter the dataframe based on the boolean index, we include it in the dataframe selector, for example, this only shows the Freshmen `Year == 'Fr'`

```
[28]: # show Freshmen
students[students['Year'] == 'Fr']
```

```
[28]:
```

	Name	GPA	Year
1	Bob	NaN	Fr
2	Chris	3.4	Fr
6	Gus	3.0	Fr

1.3.7 Part One Summary

In this part we learned the basics of `pandas`. We learned how to aggregate series, build dataframes, and use the dataframe selector.

Let's put it all together to perform a data analysis.

Challenge: Get the name and gpa of the Senior student with the lowest gpa.

Algorithm:

1. filter students to seniors only
2. find `min()` GPA from seniors
3. filter the seniors to only display the GPA that equals the min GPA.

```
[29]: seniors = students[ students['Year']=='Sr']
seniors
```

```
[29]:      Name  GPA Year
4      Ed  2.5   Sr
5  Frank  3.8   Sr
```

```
[30]: lowest_gpa_sr = seniors['GPA'].min()
lowest_gpa_sr
```

```
[30]: 2.5
```

```
[31]: person = seniors[['Name','GPA']][seniors['GPA']==lowest_gpa_sr]
person
```

```
[31]:      Name  GPA
4      Ed  2.5
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

1.4 Part Two

1.4.1 2.A Iterating over a DataFrame Or Series

Sometimes its necessary to loop over your DataFrame. For example when plotting points on on a map, you will need to build the points by iterating over the values in the dataframe.

Looping over the DataFrame To loop over a dataframe use the `iterrows()` method function. This provides an `int` index and a `dict` for the values in each column. The key in the `dict` is the column name and the value is the row value at that index.

```
[32]: students
```

```
[32]:      Name  GPA Year
0  Allen  4.0   So
1    Bob  NaN   Fr
2  Chris  3.4   Fr
3   Dave  2.8   Jr
4     Ed  2.5   Sr
5  Frank  3.8   Sr
6    Gus  3.0   Fr
```

```
[33]: for index, row in students.iterrows():
        print(f"{index}: {row['Name']} {row['GPA']} {row['Year']}")
```

```
0: Allen 4.0 So
1: Bob nan Fr
2: Chris 3.4 Fr
3: Dave 2.8 Jr
4: Ed 2.5 Sr
5: Frank 3.8 Sr
6: Gus 3.0 Fr
```

Likewise iterating over a series is similar. The `Series` itself is iterable so there is no need to use a method function.

```
[34]: for name in students['Name']:
        print(name)
```

```
Allen
Bob
Chris
Dave
Ed
Frank
Gus
```

1.4.2 2.B Loading JSON with Pandas

The Pandas `read_json()` method can easily read in list-oriented JSON-formatted text into a `DataFrame`. Each item in the list becomes a row in the `DataFrame`.

For example, here's the `students.json` file we read in with Pandas. Notice the JSON starts with a list [

```
[
  {"Name": "Abby", "GPA": 4.0, "Year": "So"},
  {"Name": "Bette", "GPA": 3.7, "Year": "Jr"},
  {"Name": "Chris", "Year": "Fr"},
  {"Name": "Dee", "GPA": 3.4, "Year": "Fr"}
]
```

```
[35]: students = pd.read_json("https://raw.githubusercontent.com/mafudge/datasets/
    ↪master/json-samples/students.json")
students
```

```
[35]:
```

	Name	GPA	Year
0	Abby	4.0	So
1	Bette	3.7	Jr
2	Chris	NaN	Fr
3	Dee	3.4	Fr

json_normalize() for nested JSON The `read_json()` method does not perform well on nested JSON structures. For example consider the following JSON file of customer orders:

The file `orders.json`:

```
[
  {
    "Customer" : { "FirstName" : "Abby", "LastName" : "Kuss"},
    "Items" : [
      { "Name" : "T-Shirt", "Price" : 10.0, "Quantity" : 3},
      { "Name" : "Jacket", "Price" : 20.0, "Quantity" : 1}
    ]
  },
  {
    "Customer" : { "FirstName" : "Bette", "LastName" : "Alott"},
    "Items" : [
      { "Name" : "Shoes", "Price" : 25.0, "Quantity" : 1},
      { "Name" : "Jacket", "Price" : 20.0, "Quantity" : 1}
    ]
  },
  {
    "Customer" : { "FirstName" : "Chris", "LastName" : "Peanugget"},
    "Items" : [
      { "Name" : "T-Shirt", "Price" : 10.0, "Quantity" : 1}
    ]
  }
]
```

When we read this with `read_json_()` we get the three orders but only two columns.

```
[36]: orders = pd.read_json("https://raw.githubusercontent.com/mafudge/datasets/
↳master/json-samples/orders.json")
orders
```

```
[36]:
```

	Customer \
0	{'FirstName': 'Abby', 'LastName': 'Kuss'}
1	{'FirstName': 'Bette', 'LastName': 'Alott'}
2	{'FirstName': 'Chris', 'LastName': 'Peanugget'}

	Items
0	[{'Name': 'T-Shirt', 'Price': 10.0, 'Quantity': ...
1	[{'Name': 'Shoes', 'Price': 25.0, 'Quantity': ...
2	[{'Name': 'T-Shirt', 'Price': 10.0, 'Quantity': ...

What we want is one row per item on the the order and the customer name to be in separate columns. the `json_normalize()` method can help here.

It is important to note that `json_normalize()` does not take a file as input, but rather de-serialized json.

```
[37]: import json
with open("orders.json") as f:
    json_data = json.load(f) #de-serialize

orders = pd.json_normalize(json_data)
orders
```

```
[37]:
```

	Items	Customer.FirstName	\
0	[{'Name': 'T-Shirt', 'Price': 10.0, 'Quantity': ...	Abby	
1	[{'Name': 'Shoes', 'Price': 25.0, 'Quantity': ...	Bette	
2	[{'Name': 'T-Shirt', 'Price': 10.0, 'Quantity': ...	Chris	

	Customer.LastName
0	Kuss
1	Alott
2	Peanugget

Better but we still need to handle the list of `Items`. To accomplish this we

1. set the `record_path` to be the nested list `'Items'`. This tells `json_normalize()` to use that JSON key as the row level. So now we will have 5 rows (one for each item) instead of 3.
2. Then we set the `meta` named argument to a list of each of the other values we wish to include, in this instance last name and first name.

NOTE: The `meta` syntax is a bit weird. Its a list of JSON paths (also represented as lists) to each item in the JSON. For example:

The meta Argument ==> Matches This in the JSON ==> And Displays As This Pandas

`["Customer", "FirstName"] ==> { "Customer" : { "FirstName": ...} ==> Customer.Firstname`

```
[38]: orders = pd.json_normalize(json_data, record_path="Items",
    ↪meta=[["Customer", "FirstName"], ["Customer", "LastName"]])
orders
```

```
[38]:
```

	Name	Price	Quantity	Customer.FirstName	Customer.LastName
0	T-Shirt	10.0	3	Abby	Kuss
1	Jacket	20.0	1	Abby	Kuss
2	Shoes	25.0	1	Bette	Alott
3	Jacket	20.0	1	Bette	Alott
4	T-Shirt	10.0	1	Chris	Peanugget

Yes it seems complicated, because it is a bit complicated, let's try another example, with some abstract values.

In the following example we want to generate a normalized table with 3 rows and 4 columns.

- The rows are based on the “A” `record_path`. There are three: 101, 111 and 201
- The meta data are based on columns “B”, and “C1”

```
[39]: json_data = [
      {
        "A": [
          {"A1": 101, "A2": 102},
          {"A1": 111, "A2": 112}
        ],
        "B": 103,
        "C": {"C1": 104}
      },
      {
        "A": [
          {"A1": 201, "A2": 202}
        ],
        "B": 203,
        "C": {"C1": 204}
      }
    ]
```

```
[40]: df = pd.json_normalize(json_data, record_path="A", meta=["B", ["C", "C1"]])
df
```

```
[40]:
```

	A1	A2	B	C.C1
0	101	102	103	104
1	111	112	103	104
2	201	202	203	204

1.4.3 2.C Simple Web scraping with Pandas

The Pandas `read_html()` method can be used to extract HTML tables from websites. For any given webpage url as input, `read_html()` will output a list of DataFrames. There will be one DataFrame for each table on the website.

In this example we read the following wikipedia page: https://en.wikipedia.org/wiki/National_Basketball_Association

NOTE: We use wayback machine to snapshot the page to the point in time this paper was written.

```
[41]: dfs = pd.read_html("https://web.archive.org/web/20240316160100/https://en.
↳wikipedia.org/wiki/National_Basketball_Association")
```

Upon inspection the NBA teams are the 4th table in the list.

```
[42]: teams = dfs[3]
teams.head(n=5)
```

```
[42]:
```

	Conference	Division	Team \
0	Eastern Conference	Atlantic	Boston Celtics
1	Eastern Conference	Atlantic	Brooklyn Nets
2	Eastern Conference	Atlantic	New York Knicks
3	Eastern Conference	Atlantic	Philadelphia 76ers

4	Eastern Conference	Atlantic	Toronto Raptors	
---	--------------------	----------	-----------------	--

	Location	Arena	Capacity	\
0	Boston, Massachusetts	TD Garden	19156	
1	New York, New York	Barclays Center	17732	
2	New York, New York	Madison Square Garden	19812	
3	Philadelphia, Pennsylvania	Wells Fargo Center	20478	
4	Toronto, Ontario	Scotiabank Arena	19800	

	Coordinates	Founded	Joined
0	42°21 59 N 71°03 44 W / 42.366303°N 71.062228°W	1946	1946
1	40°40 58 N 73°58 29 W / 40.68265°N 73.974689°W	1967*	1976
2	40°45 02 N 73°59 37 W / 40.750556°N 73.993611°W	1946	1946
3	39°54 04 N 75°10 19 W / 39.901111°N 75.171944°W	1946*	1949
4	43°38 36 N 79°22 45 W / 43.643333°N 79.379167°W	1995	1995

A table of championships can be found as the 6th table.

```
[43]: championships = dfs[5]
      championships.head(n=5)
```

```
[43]:
```

	Teams	Win	Loss	Total	\
0	Minneapolis/Los Angeles Lakers	17	15	32	
1	Boston Celtics	17	5	22	
2	Philadelphia/San Francisco/Golden State Warriors	7	5	12	
3	Chicago Bulls	6	0	6	
4	San Antonio Spurs	5	1	6	

	Year(s) won	\
0	1949, 1950, 1952, 1953, 1954, 1972, 1980, 1982...	
1	1957, 1959, 1960, 1961, 1962, 1963, 1964, 1965...	
2	1947, 1956, 1975, 2015, 2017, 2018, 2022	
3	1991, 1992, 1993, 1996, 1997, 1998	
4	1999, 2003, 2005, 2007, 2014	

	Year(s) runner-up
0	1959, 1962, 1963, 1965, 1966, 1968, 1969, 1970...
1	1958, 1985, 1987, 2010, 2022
2	1948, 1964, 1967, 2016, 2019
3	-
4	2013

You'll have to investigate the output to find what you need There is no magic command you can type to find the table you want. Ultimately you will need to inspect each DataFrame returned from `read_html()` and figure out which DataFrame in the list contains the table of data you need.

Here's some sample code to list every table on the webpage along with its index. We simple use

a for loop to iterate over each DataFrame. Code like this can help you to identify the table you want from the webpage.

```
[44]: from IPython.display import display

webpage = "https://web.archive.org/web/20230606035218/https://ist256.com/"
dataframes = pd.read_html(webpage)
for index, df in enumerate(dataframes):
    print("INDEX:", index)
    display(df.head(5))
```

INDEX: 0

	0	1	2
0	Feb	JUN	Dec
1	NaN	06	NaN
2	2022	2023	2024

INDEX: 1

	Dates	Topic (Click Link for Content and Assigned Readings)
0	1/18 - 1/22	Wednesday is the first class.
1	1/23 - 1/29	Lesson 01: Introduction to Python Programming
2	1/30 - 2/5	Lesson 02: Input, Output, Variables and Types
3	2/6 - 2/12	Lesson 03: Conditionals
4	2/13 - 2/19	Lesson 04: Iterations

INDEX: 2

	Date Due	Time Due	Gradebook	Points	Tool \
0	1/24/2023	11:59 PM	L01	3	Jupyterhub
1	1/25/2023	11:59 PM	S01	3	Jupyterhub
2	1/27/2023	11:59 PM	H01	3	Jupyterhub
3	1/31/2023	11:59 PM	L02	3	Jupyterhub
4	2/1/2023	11:59 PM	S02	3	Jupyterhub

	What is Due?
0	01-Intro/LAB-Intro.ipynb
1	01-Intro/SmallGroup-Intro.ipynb
2	01-Intro/HW-Intro.ipynb
3	02-Variables/LAB-Variables.ipynb
4	02-Variables/SmallGroup-Variables.ipynb

1.4.4 2.D Advanced Dataframe Operations

In this section we explore the various options for combining one or more dataframes together. We'll cover the two most common method functions:

- `concat()` - Row-oriented. Combine one or more DataFrames together. Appropriate for similar sets of data. Online students + Campus students.

- `merge()` - Column-oriented. Combine one or more DataFrames based on a matching column. Appropriate for different sets of data that share a business rule. Students *take* Classes so we merge students + classes so we know who is taking which class.

Concat() Example. The common use case for `concat()` is to combine two separate but similar data sets into a single data set.

```
[45]: campus_students = pd.read_csv("https://raw.githubusercontent.com/mafudge/
↳datasets/master/delimited/campus-students.csv")
campus_students
```

```
[45]:
```

	Name	Grade	Year
0	Helen	NaN	Sophomore
1	Iris	10.0	Senior
2	Jimmy	8.0	Freshman
3	Karen	NaN	Freshman
4	Lynne	10.0	Sophomore
5	Mike	10.0	Sophomore
6	Nico	NaN	Junior
7	Pete	8.0	Freshman

```
[46]: online_students = pd.read_csv("https://raw.githubusercontent.com/mafudge/
↳datasets/master/delimited/online-students.csv")
online_students
```

```
[46]:
```

	Name	Grade	Year	Location
0	Abby	7.0	Freshman	NY
1	Bob	9.0	Sophomore	CA
2	Chris	10.0	Senior	CA
3	Dave	8.0	Freshman	NY
4	Ellen	7.0	Sophomore	TX
5	Fran	10.0	Senior	FL
6	Greg	8.0	Freshman	NY

To perform the `concat()` we provide a list of DataFrames to concatenate.

```
[47]: students = pd.concat([campus_students, online_students])
students
```

```
[47]:
```

	Name	Grade	Year	Location
0	Helen	NaN	Sophomore	NaN
1	Iris	10.0	Senior	NaN
2	Jimmy	8.0	Freshman	NaN
3	Karen	NaN	Freshman	NaN
4	Lynne	10.0	Sophomore	NaN
5	Mike	10.0	Sophomore	NaN
6	Nico	NaN	Junior	NaN
7	Pete	8.0	Freshman	NaN

0	Abby	7.0	Freshman	NY
1	Bob	9.0	Sophomore	CA
2	Chris	10.0	Senior	CA
3	Dave	8.0	Freshman	NY
4	Ellen	7.0	Sophomore	TX
5	Fran	10.0	Senior	FL
6	Greg	8.0	Freshman	NY

concat() - Ignoring the index As you can see from the code above the index from the original DataFrames was used. For example Helen and Abby both share the index 0. While this is acceptable, there are situations where a new index based on combined values is desirable. To make this happen include the `ignore_index=True` named argument. This will create a new index from the output DataFrame.

```
[48]: students = pd.concat([campus_students, online_students], ignore_index=True)
students
```

```
[48]:
```

	Name	Grade	Year	Location
0	Helen	NaN	Sophomore	NaN
1	Iris	10.0	Senior	NaN
2	Jimmy	8.0	Freshman	NaN
3	Karen	NaN	Freshman	NaN
4	Lynne	10.0	Sophomore	NaN
5	Mike	10.0	Sophomore	NaN
6	Nico	NaN	Junior	NaN
7	Pete	8.0	Freshman	NaN
8	Abby	7.0	Freshman	NY
9	Bob	9.0	Sophomore	CA
10	Chris	10.0	Senior	CA
11	Dave	8.0	Freshman	NY
12	Ellen	7.0	Sophomore	TX
13	Fran	10.0	Senior	FL
14	Greg	8.0	Freshman	NY

merge() Example. The common use case for `merge()` is to combine two separate but dissimilar data sets into a single data set. These DataFrames **share a common column** so you can match the output of each DataFrame to the other on matching rows.

```
[49]: customers = pd.read_csv("https://raw.githubusercontent.com/mafudge/datasets/
↳master/delimited/customers.csv")
customers
```

```
[49]:
```

	customer_id	firstname	lastname
0	10	Abby	Kuss
1	20	Bette	Alott
2	30	Chris	Peanugget

```
[50]: orders = pd.read_csv("https://raw.githubusercontent.com/mafudge/datasets/master/
↳ delimited/orders.csv")
orders
```

```
[50]:   order_id  order_customer_id  item  price  qty
0      1001             10  T-Shirt    10    3
1      1002             10   Jacket    20    1
2      1003             20   Shoes    25    1
3      1004             20   Jacket    20    1
4      1005             30  T-Shirt    10    1
```

For the `customers` and `orders` DataFrames we want to combine them together so that each customer is matched to their order. In this case the common column in both DataFrames is `customer_id` (from `customers`) and `order_customer_id` (from `orders`).

To perform the merge we need at least 4 named arguments are as follows:

- `left` the DataFrame on the left.
- `right` the DataFrame on the right.
- `left_on` - the column from the `left` dataframe to match
- `right_on` - the column from the `right` dataframe to match

```
[51]: combined = pd.merge(left=customers, right=orders, left_on="customer_id",
↳ right_on="order_customer_id")
combined
```

```
[51]:   customer_id  firstname  lastname  order_id  order_customer_id  item  \
0           10        Abby        Kuss      1001             10  T-Shirt
1           10        Abby        Kuss      1002             10   Jacket
2           20        Bette        Alott      1003             20   Shoes
3           20        Bette        Alott      1004             20   Jacket
4           30        Chris  Peanugget      1005             30  T-Shirt

   price  qty
0      10    3
1      20    1
2      25    1
3      20    1
4      10    1
```

Left and Right are relative Of course left and right are relative, as we need to know which dataframe we are referring to as part of the merge. For example this is the same merge:

```
[52]: combined = pd.merge(left=orders, right=customers, left_on="order_customer_id",
↳ right_on="customer_id")
combined
```

```
[52]:
```

	order_id	order_customer_id	item	price	qty	customer_id	firstname \
0	1001	10	T-Shirt	10	3	10	Abby
1	1002	10	Jacket	20	1	10	Abby
2	1003	20	Shoes	25	1	20	Bette
3	1004	20	Jacket	20	1	20	Bette
4	1005	30	T-Shirt	10	1	30	Chris

	lastname
0	Kuss
1	Kuss
2	Alott
3	Alott
4	Peanugget

1.4.5 2.E Generating Columns, Lambdas and the apply() function

In this section we demonstrate how to derive new columns in the DataFrame from the existing data.

```
[53]: data_dict = {
        'Name': ['Allen', 'Bob', 'Chris', 'Dave', 'Ed', 'Frank', 'Gus'],
        'GPA': [4.0, np.nan, 3.4, 2.8, 2.5, 3.8, 3.0],
        'Year': ['So', 'Fr', 'Fr', 'Jr', 'Sr', 'Sr', 'Fr'] }
students = pd.DataFrame( data_dict )
students
```

```
[53]:
```

	Name	GPA	Year
0	Allen	4.0	So
1	Bob	NaN	Fr
2	Chris	3.4	Fr
3	Dave	2.8	Jr
4	Ed	2.5	Sr
5	Frank	3.8	Sr
6	Gus	3.0	Fr

Column generation Pandas allows us to create new columns from existing data. This is simple to do when the values are constants or simple formulas.

This adds a column "MaxGPA" with the set value of 4.0 for every value in the series.

```
[54]: students["MaxGPA"] = 4.0
students
```

```
[54]:
```

	Name	GPA	Year	MaxGPA
0	Allen	4.0	So	4.0
1	Bob	NaN	Fr	4.0
2	Chris	3.4	Fr	4.0
3	Dave	2.8	Jr	4.0

4	Ed	2.5	Sr	4.0
5	Frank	3.8	Sr	4.0
6	Gus	3.0	Fr	4.0

This creates a calculated column "DiffGPA" which is the difference between two columns:

```
[55]: students["DiffGPA"] = students["MaxGPA"] - students["GPA"]
students
```

```
[55]:
```

	Name	GPA	Year	MaxGPA	DiffGPA
0	Allen	4.0	So	4.0	0.0
1	Bob	NaN	Fr	4.0	NaN
2	Chris	3.4	Fr	4.0	0.6
3	Dave	2.8	Jr	4.0	1.2
4	Ed	2.5	Sr	4.0	1.5
5	Frank	3.8	Sr	4.0	0.2
6	Gus	3.0	Fr	4.0	1.0

Lambda Functions For more complex logic, a user-defined function can be created and then executed for each value using the `apply()` method function.

For example, consider the following function code to calculate whether an input GPA is on the deans list ≥ 3.4 :

```
[56]: def deans_list(gpa: float) -> str:
        if gpa >= 3.4:
            return "Yes"
        else:
            return "No"

deans_list(3.4)
```

```
[56]: 'Yes'
```

Here's how we use the function with the `apply()` method to create a new Pandas column, breaking it down:

- `apply()` operates on a DataFrame, in this case `students`.
- `apply()` returns a Series, which we add to the column "DeansList"
- `lambda row:` is a way of saying "do this for each row"
- `row` becomes an iterator of `dict` of values in each row. For example in the first row (`index=0`) `row['Name'] == 'Allen'` and `row['GPA'] == 4.0`
- `axis=1` tells `apply()` to loop over rows. This is almost always what we want.

```
[57]: students['DeansList'] = students.apply(lambda row: deans_list(row['GPA']),
        ↪axis=1)
students
```

```
[57]:
```

	Name	GPA	Year	MaxGPA	DiffGPA	DeansList
0	Allen	4.0	So	4.0	0.0	Yes
1	Bob	NaN	Fr	4.0	NaN	No
2	Chris	3.4	Fr	4.0	0.6	Yes
3	Dave	2.8	Jr	4.0	1.2	No
4	Ed	2.5	Sr	4.0	1.5	No
5	Frank	3.8	Sr	4.0	0.2	Yes
6	Gus	3.0	Fr	4.0	1.0	No

1.4.6 2.F DataFrames as output

Throughout most of this tutorial we have been echoing the DataFrame content to the Jupyter cell. In this section we discuss your options for rendering the DataFrame as output when part of a larger program.

```
[58]: data_dict = {
        'Name':  ['Allen', 'Bob', 'Chris', 'Dave', 'Ed', 'Frank', 'Gus'],
        'GPA':   [4.0, np.nan, 3.4, 2.8, 2.5, 3.8, 3.0],
        'Year' : ['So', 'Fr', 'Fr', 'Jr', 'Sr', 'Sr', 'Fr'] }
students = pd.DataFrame( data_dict )
```

Printing To output a dataframe in code, you can `print()` it. You'll get a text-based version of the dataframe. For example:

```
[59]: size = int(input("How many students to display at random? "))
random_students = students.sample(n=size)
print(random_students)
```

How many students to display at random? 4

```

      Name  GPA  Year
1    Bob  NaN   Fr
5  Frank  3.8   Sr
2  Chris  3.4   Fr
3   Dave  2.8   Jr
```

IPython.display To get a well-formatted dataframe you can pass it in as an argument to `IPython.display`. The same example:

```
[60]: from IPython.display import display
from ipywidgets import interact_manual
max_students = len(students)
@interact_manual(size=(1,max_students,1))
def on_click(size):
    random_students = students.sample(n=size)
    display(random_students)
```



```
interactive(children=(IntSlider(value=4, description='size', max=7, min=1),  
    Button(description='Run Interact',...
```

[]:

[]:

[]: