

Practice-PandasTitanic

October 7, 2024

1 Name: Ramya Chowdary Patchala

2 Data Analysis with Pandas - PRACTICE 1

We will perform a data analysis on the **RMS Titanic** passenger list. The RMS Titanic is one of the most famous ocean liners in history. On April 15, 1912 it sank after colliding with an iceberg in the North Atlantic Ocean. To learn more, read here: https://en.wikipedia.org/wiki/RMS_Titanic

Our goal today is to perform a data analysis on a subset of the passenger list. We're looking for insights as to which types of passengers did and didn't survive. Women? Children? 1st Class Passengers? 3rd class? Etc.

I'm sure you've heard the expression often said during emergencies: "Women and Children first" Let's explore this data set and find out if that's true!

Before we begin you should read up on what each of the columns mean in the data dictionary. You can find this information on this page: <https://www.kaggle.com/c/titanic/data>

2.1 Loading the data set

First we load the dataset into a Pandas `DataFrame` variable. The `sample(10)` method takes a random sample of 10 passengers from the data set.

```
[1]: import pandas as pd
import numpy as np

# this turns off warning messages
import warnings
warnings.filterwarnings('ignore')

passengers = pd.read_csv('https://raw.githubusercontent.com/mafudge/datasets/
↪master/ist256/12-pandas/titanic.csv')
passengers.sample(10)
```

```
[1]:
```

	PassengerId	Survived	Pclass	Name \
340	341	1	2	Navratil, Master. Edmond Roger
306	307	1	1	Fleming, Miss. Margaret
153	154	0	3	van Billiard, Mr. Austin Blyler
853	854	1	1	Lines, Miss. Mary Conover
622	623	1	3	Nakid, Mr. Sahid

6	7	0	1	McCarthy, Mr. Timothy J
82	83	1	3	McDermott, Miss. Brigdet Delia
775	776	0	3	Myhrman, Mr. Pehr Fabian Oliver Malkolm
816	817	0	3	Heininen, Miss. Wendla Maria
317	318	0	2	Moraweck, Dr. Ernest

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
340	male	2.0	1	1	230080	26.0000	F2	S
306	female	NaN	0	0	17421	110.8833	NaN	C
153	male	40.5	0	2	A/5. 851	14.5000	NaN	S
853	female	16.0	0	1	PC 17592	39.4000	D28	S
622	male	20.0	1	1	2653	15.7417	NaN	C
6	male	54.0	0	0	17463	51.8625	E46	S
82	female	NaN	0	0	330932	7.7875	NaN	Q
775	male	18.0	0	0	347078	7.7500	NaN	S
816	female	23.0	0	0	STON/O2. 3101290	7.9250	NaN	S
317	male	54.0	0	0	29011	14.0000	NaN	S

2.2 How many survived?

One of the first things we should do is figure out how many of the passengers in this data set survived. Let's start with isolating just the 'Survived' column into a series:

```
[2]: passengers
```

```
[2]: PassengerId  Survived  Pclass  \
0              1         0       3
1              2         1       1
2              3         1       3
3              4         1       1
4              5         0       3
..          ...         ...     ...
886           887         0       2
887           888         1       1
888           889         0       3
889           890         1       1
890           891         0       3
```

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	
..	
886	Montvila, Rev. Juozas	male	27.0	0	
887	Graham, Miss. Margaret Edith	female	19.0	0	

888	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1
889	Behr, Mr. Karl Howell	male	26.0	0
890	Dooley, Mr. Patrick	male	32.0	0

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
..
886	0	211536	13.0000	NaN	S
887	0	112053	30.0000	B42	S
888	2	W./C. 6607	23.4500	NaN	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	NaN	Q

[891 rows x 12 columns]

```
[3]: passengers['Survived'].sample(10)
```

```
[3]: 631    0
      34    0
      343  0
      869  1
      870  0
      452  0
      459  0
      471  0
      377  0
      554  1
      Name: Survived, dtype: int64
```

There's too many to display so we just display a random sample of 10 passengers.

- 1 means the passenger survived
- 0 means the passenger died

What we really want is to count the number of survivors and deaths. We do this by querying the `value_counts()` of the `['Survived']` column, which returns a `Series` of counts, like this:

```
[4]: passengers['Survived'].value_counts()
```

```
[4]: Survived
      0    549
      1    342
      Name: count, dtype: int64
```

Only 342 passengers survived, and 549 perished. Let's observe this same data as percentages of

the whole. We do this by adding the `normalize=True` named argument to the `value_counts()` method.

```
[5]: passengers['Survived'].value_counts(normalize=True)
```

```
[5]: Survived
0    0.616162
1    0.383838
Name: proportion, dtype: float64
```

Just 38% of passengers in this dataset survived.

2.2.1 1.1 You Code

FIRST Write a Pandas expression to display counts of males and female passengers using the `Sex` variable:

```
[6]: # todo write code here
passengers["Sex"].value_counts()
```

```
[6]: Sex
male      577
female    314
Name: count, dtype: int64
```

2.2.2 1.2 You Code

NEXT Write a Pandas expression to display male /female passenger counts as a percentage of the whole number of passengers in the data set.

```
[13]: # todo write code here
passengers["Sex"].value_counts(normalize=True)
```

```
[13]: Sex
male      0.647587
female    0.352413
Name: proportion, dtype: float64
```

If you got things working, you now know that **35% of passengers were female**.

2.3 Who survives? Men or Women?

We now know that 35% of the passengers were female, and 65% we male.

The next thing to think about is how do survival rates affect these numbers?

If the ratio is about the same for survivors only, then we can conclude that your **Sex** did not play a role in your survival on the RMS Titanic.

Let's find out.

```
[14]: survivors = passengers[passengers['Survived'] ==1]
survivors['PassengerId'].count()
```

```
[14]: 342
```

Still **342** like we discovered originally. Now let's check the **Sex** split among survivors only:

```
[15]: survivors['Sex'].value_counts()
```

```
[15]: Sex
female    233
male      109
Name: count, dtype: int64
```

WOW! That is a huge difference! But you probably can't see it easily. Let's represent it in a DataFrame, so that it's easier to visualize:

```
[16]: sex_all_series = passengers['Sex'].value_counts()
sex_survivor_series = survivors['Sex'].value_counts()

sex_comparision_df = pd.DataFrame({ 'AllPassengers' : sex_all_series,
    ↪ 'Survivors' : sex_survivor_series })
sex_comparision_df['SexSurvivialRate'] = sex_comparision_df['Survivors'] /
    ↪ sex_comparision_df['AllPassengers']
sex_comparision_df
```

```
[16]:
```

	AllPassengers	Survivors	SexSurvivialRate
Sex			
female	314	233	0.742038
male	577	109	0.188908

```
[17]: sex_all_series = passengers['Sex'].value_counts()
sex_all_series
```

```
[17]: Sex
male      577
female    314
Name: count, dtype: int64
```

So, females had a 74% survival rate. Much better than the overall rate of 38%

We should probably briefly explain the code above.

- The first two lines get a series count of all passengers by Sex (male / female) and count of survivors by sex
- The third line creates a Pandas DataFrame. Recall a pandas dataframe is just a dictionary of series. We have two keys 'AllPassengers' and 'Survivors'
- The fourth line creates a new column in the dataframe which is just the survivors / all passengers to get the rate of survival for that Sex.

2.4 Feature Engineering: Adults and Children

Sometimes the variable we want to analyze is not readily available, but can be created from existing data. This is commonly referred to as **feature engineering**. The name comes from machine learning where we use data called *features* to predict an outcome.

Let's create a new feature called 'AgeCat' as follows:

- When **Age** ≤ 18 then 'Child'
- When **Age** > 18 then 'Adult'

This is easy to do in pandas. First we create the column and set all values to `np.nan` which means 'Not a number'. This is Pandas way of saying no value. Then we set the values based on the rules we set for the feature.

```
[18]: passengers['AgeCat'] = np.nan # Not a number
passengers['AgeCat'][ passengers['Age'] <=18 ] = 'Child'
passengers['AgeCat'][ passengers['Age'] > 18 ] = 'Adult'
passengers.sample(5)
```

```
[18]:
```

	PassengerId	Survived	Pclass	Name	Sex	\
701	702	1	1	Silverthorne, Mr. Spencer Victor	male	
590	591	0	3	Rintamaki, Mr. Matti	male	
750	751	1	2	Wells, Miss. Joan	female	
598	599	0	3	Boulos, Mr. Hanna	male	
61	62	1	1	Icard, Miss. Amelie	female	

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	AgeCat
701	35.0	0	0	PC 17475	26.2875	E24	S	Adult
590	35.0	0	0	STON/O 2. 3101273	7.1250	NaN	S	Adult
750	4.0	1	1	29103	23.0000	NaN	S	Child
598	NaN	0	0	2664	7.2250	NaN	C	NaN
61	38.0	0	0	113572	80.0000	B28	NaN	Adult

Let's get the count and distributions of Adults and Children on the passenger list.

```
[19]: passengers
```

```
[19]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	
..	
886	887	0	2	
887	888	1	1	
888	889	0	3	
889	890	1	1	
890	891	0	3	

		Name	Sex	Age	SibSp	\
0		Braund, Mr. Owen Harris	male	22.0	1	
1	Cummings, Mrs. John Bradley (Florence Briggs Th...		female	38.0	1	
2		Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)		female	35.0	1	
4		Allen, Mr. William Henry	male	35.0	0	
..			
886		Montvila, Rev. Juozas	male	27.0	0	
887		Graham, Miss. Margaret Edith	female	19.0	0	
888	Johnston, Miss. Catherine Helen "Carrie"		female	NaN	1	
889		Behr, Mr. Karl Howell	male	26.0	0	
890		Dooley, Mr. Patrick	male	32.0	0	

	Parch	Ticket	Fare	Cabin	Embarked	AgeCat
0	0	A/5 21171	7.2500	NaN	S	Adult
1	0	PC 17599	71.2833	C85	C	Adult
2	0	STON/O2. 3101282	7.9250	NaN	S	Adult
3	0	113803	53.1000	C123	S	Adult
4	0	373450	8.0500	NaN	S	Adult
..	
886	0	211536	13.0000	NaN	S	Adult
887	0	112053	30.0000	B42	S	Adult
888	2	W./C. 6607	23.4500	NaN	S	NaN
889	0	111369	30.0000	C148	C	Adult
890	0	370376	7.7500	NaN	Q	Adult

[891 rows x 13 columns]

And here's the percentage as a whole:

```
[20]: passengers['AgeCat'].value_counts(normalize=True)
```

```
[20]: AgeCat
Adult    0.805322
Child    0.194678
Name: proportion, dtype: float64
```

So close to **80%** of the passengers were adults. Once again let's look at the ratio of **AgeCat** for survivors only. If your age has no bearing of survival, then the rates should be the same.

Here are the counts of Adult / Children among the survivors only:

```
[21]: survivors = passengers[passengers['Survived'] ==1]
survivors['AgeCat'].value_counts()
```

```
[21]: AgeCat
Adult    220
```

```
Child      70
Name: count, dtype: int64
```

2.4.1 1.3 You Code

Calculate the AgeCat survival rate, similar to how we did for the SexSurvivalRate.

```
[23]: agecat_all_series = passengers['AgeCat'].value_counts()
      agecat_survivor_series = survivors['AgeCat'].value_counts()

      # todo make a data frame, add AgeCatSurvivalRate column, display dataframe
      age_comparison_df = pd.DataFrame({'All': agecat_all_series, 'Survivors':
      ↪agecat_survivor_series })
      age_comparison_df['AgeSurvivalRate'] = age_comparison_df['Survivors']/
      ↪age_comparison_df['All']
      age_comparison_df
```

```
[23]:      All  Survivors  AgeSurvivalRate
AgeCat
Adult    575         220         0.382609
Child    139          70         0.503597
```

So, children had a 50% survival rate, better than the overall rate of 38%

2.5 So, women and children first?

It looks like the RMS really did have the motto: “Women and Children First.”

Here are our insights. We know:

- If you were a passenger, you had a 38% chance of survival.
- If you were a female passenger, you had a 74% chance of survival.
- If you were a child passenger, you had a 50% chance of survival.

2.5.1 Now you try it for Passenger Class

Repeat this process for Pclass The passenger class variable. Display the survival rates for each passenger class. What does the information tell you about passenger class and survival rates?

I'll give you a hint... “Class matters!”

2.5.2 1.4 You Code

```
[24]: # todo: repeat the analysis in the previous cell for Pclass
      all_pclass_series= passengers['Pclass'].value_counts()
      survived_pclass_series = passengers[ passengers['Survived'] == 1]['Pclass'].
      ↪value_counts()
      pclass_df = pd.DataFrame( { 'All' : all_pclass_series, 'Survived' :
      ↪survived_pclass_series})
      pclass_df['Ratio'] = pclass_df['Survived'] /pclass_df['All']
```



```
pclass_df
```

```
[24]:
```

	All	Survived	Ratio
Pclass			
1	216	136	0.629630
2	184	87	0.472826
3	491	119	0.242363

```
[25]: passengers[ passengers['Survived'] == 1]['Pclass'].value_counts()
```

```
[25]: Pclass
1      136
3      119
2       87
Name: count, dtype: int64
```

Not a big surprise. The 1st class passengers had a 62.9% survival rate!

2.6 What have we learned?

Your best odds of survival were:

- First class ticket Pclass=1
- Female
- Child

Your job is to check the survival rate of those individuals. Here's the process

1. filter the passengers data frame by the above criteria
2. normalize the value counts of survived.

Learn that while only 38% of all passengers survived, 90.9% passengers meeting this criteria survived!

2.6.1 1.5 You Code

```
[26]: # TODO
ffc_df = passengers[passengers['Pclass'] == 1][passengers['Sex']=='
↳ 'female'][passengers['AgeCat']=='Child']
ffc_df['Survived'].value_counts(normalize = True)
```

```
[26]: Survived
1      0.909091
0      0.090909
Name: proportion, dtype: float64
```

3 Metacognition

3.0.1 Rate your comfort level with this week's material so far.

1 ==> I don't understand this at all yet and need extra help. If you choose this please try to articulate that which you do not understand to the best of your ability in the questions and comments section below.

2 ==> I can do this with help or guidance from other people or resources. If you choose this level, please indicate HOW this person helped you in the questions and comments section below.

3 ==> I can do this on my own without any help.

4 ==> I can do this on my own and can explain/teach how to do it to others.

---= Double-Click Here then Enter a Number 1 through 4 Below This Line ==--

3.0.2 4